

Project 2: TMDb Movie Data Analysis

Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
- [Conclusions](#)

Introduction

Questions to Explore:

- [Q0: Fun Facts](#)
- [Q1: Genre Trends from 1960 to 2015](#)
- [Q2: Properties Associated With Higher Profits](#)

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set()
```

Data Wrangling

General Properties


In [2]: *# Load data and print out a few lines.*

```
df = pd.read_csv("tmdb_movies.csv")
df.head()
```

Out[2]:

	id	imdb_id	popularity	budget	revenue	original_title	cast	
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	http://w
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D...	I
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	Vin Diesel Paul Walker Jason Statham Michelle ...	

5 rows × 21 columns



In [3]: *# Assess number of rows and columns of dataset*

```
df.shape
```

Out[3]: (10866, 21)

In [4]: *# Assess dataset, including datatypes, and check for missing data.*
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count        10866 non-null int64
vote_average      10866 non-null float64
release_year      10866 non-null int64
budget_adj        10866 non-null float64
revenue_adj       10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

- Many columns have missing row data.
- First select columns for analysis and drop non-useful columns and then deal with missing data.

In [5]: df.describe()

Out[5]:

	id	popularity	budget	revenue	runtime	vote_count	vc
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000	10866.000000	10
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863	217.389748	
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405	575.619058	
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000	10.000000	
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000	17.000000	
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000	38.000000	
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000	145.750000	
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	

- There are too much information and will drop some columns when perform data analysis.
- Could add profit = (revenue - budget) and profit_ratio = (profit/budget) columns to investigate profitability
- Columns 'cast', 'director', 'keywords', 'genres', 'production_companies' contain multiple values separated by pipe (|) that need to be separated out.

Data Cleaning

First I'll remove extraneous columns that aren't relevant to my analysis and duplicates rows. Then I'll add and/or replace information to ensure my dataset is clean for analysis.

Need to drop columns:

- 'imdb_id' : Already have 'id'
- 'homepage' : Not relevant
- 'tagline' : Not relevant
- 'overview' : Not relevant
- 'release_year' : Already have 'release_date'

```
In [6]: # Drop Columns
df.drop(['imdb_id', 'homepage', 'tagline', 'overview', 'release_date'], axis = 1,
inplace = True)
df.head(1)
```

Out[6]:

	id	popularity	budget	revenue	original_title	cast	director	
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	monster d rex v

- Add 'profit' and 'profit_adj' columns

```
In [7]: df['profit'] = df['revenue'] - df['budget']
df['profit_adj'] = df['revenue_adj'] - df['budget_adj']
```

- Add 'profit_ratio' and 'profit_ratio_adj' columns
- Add 0.000001 to revenue to prevent NaN in ratios

```
In [8]: df['profit_ratio'] = df['profit']/(0.000001+df['budget'])
df['profit_ratio_adj'] = df['profit_adj']/(0.000001+df['budget_adj'])
```

Check/Drop for duplicates

```
In [9]: df[df.duplicated()]
```

Out[9]:

	id	popularity	budget	revenue	original_title	cast	director	keyw
2090	42194	0.59643	30000000	967000	TEKKEN	Jon Foo Kelly Overton Cary-Hiro... Tagawa lan...	Dwight H. Little	arts dystopia b on \ game nr

```
In [10]: df.query('original_title == "TEKKEN"')
```

Out[10]:

	id	popularity	budget	revenue	original_title	cast	director	keyw
2089	42194	0.59643	30000000	967000	TEKKEN	Jon Foo Kelly Overton Cary-Hiro... Tagawa lan...	Dwight H. Little	arts dystopia b on \ game nr
2090	42194	0.59643	30000000	967000	TEKKEN	Jon Foo Kelly Overton Cary-Hiro... Tagawa lan...	Dwight H. Little	arts dystopia b on \ game nr

```
In [11]: df.drop_duplicates(inplace = True)
print(sum(df.duplicated()))
print(df.shape)
```

```
0
(10865, 20)
```

Check for missing values

In [12]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 20 columns):
id                10865 non-null int64
popularity        10865 non-null float64
budget            10865 non-null int64
revenue           10865 non-null int64
original_title    10865 non-null object
cast              10789 non-null object
director          10821 non-null object
keywords          9372 non-null object
runtime           10865 non-null int64
genres            10842 non-null object
production_companies 9835 non-null object
vote_count        10865 non-null int64
vote_average      10865 non-null float64
release_year      10865 non-null int64
budget_adj        10865 non-null float64
revenue_adj       10865 non-null float64
profit            10865 non-null int64
profit_adj        10865 non-null float64
profit_ratio      10865 non-null float64
profit_ratio_adj  10865 non-null float64
dtypes: float64(7), int64(7), object(6)
memory usage: 1.7+ MB
```

In [13]: df.isnull().sum()

```
Out[13]: id                0
popularity              0
budget                  0
revenue                 0
original_title          0
cast                    76
director                44
keywords               1493
runtime                 0
genres                  23
production_companies   1030
vote_count              0
vote_average            0
release_year            0
budget_adj              0
revenue_adj             0
profit                  0
profit_adj              0
profit_ratio            0
profit_ratio_adj        0
dtype: int64
```

Drop the rows with missing values

- First drop rows with missing 'cast', 'director' and 'genres' informations

```
In [14]: df.dropna(subset = ['cast', 'director', 'genres'], inplace = True)
df.isnull().sum()
```

```
Out[14]: id                0
popularity                0
budget                  0
revenue                 0
original_title           0
cast                   0
director                0
keywords               1425
runtime                 0
genres                 0
production_companies    959
vote_count              0
vote_average            0
release_year            0
budget_adj              0
revenue_adj             0
profit                  0
profit_adj              0
profit_ratio            0
profit_ratio_adj        0
dtype: int64
```

NOTE:

df is for answering general questions.

Also need separate df_keywords, df_production, df_cast and df_director.

```
In [15]: df_keywords = df.copy()
df_keywords.dropna(subset = ['keywords'], inplace = True)
df_keywords.isnull().sum()
```

```
Out[15]: id                0
popularity                0
budget                   0
revenue                  0
original_title            0
cast                     0
director                 0
keywords                 0
runtime                  0
genres                   0
production_companies    640
vote_count               0
vote_average             0
release_year             0
budget_adj               0
revenue_adj              0
profit                   0
profit_adj               0
profit_ratio             0
profit_ratio_adj         0
dtype: int64
```

```
In [16]: df_production = df.copy()
df_production.dropna(subset = ['production_companies'], inplace = True)
df_production.isnull().sum()
```

```
Out[16]: id                0
popularity                0
budget                   0
revenue                  0
original_title            0
cast                     0
director                 0
keywords                1106
runtime                  0
genres                   0
production_companies      0
vote_count               0
vote_average             0
release_year             0
budget_adj               0
revenue_adj              0
profit                   0
profit_adj               0
profit_ratio             0
profit_ratio_adj         0
dtype: int64
```


Now have df is for answering general questions not related to keywords and production companies. Also need to make separate df_keywords and df_production for keywords and production company related questions.

Split up 'genres' columns

```
In [17]: df_split_genre = df.copy()
split_genre = df_split_genre['genres'].str.split('|').apply(pd.Series,1).stack()
().reset_index(level=1, drop=True)
split_genre.name = 'genre_split'
df_split_genre = df_split_genre.drop(['genres'], axis=1).join(split_genre)
df_split_genre.head(3)
```

Out[17]:

	id	popularity	budget	revenue	original_title	cast	director	
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	monster drex v
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	monster drex v
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	monster drex v

In [18]: `df_split_genre.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 26753 entries, 0 to 10865
Data columns (total 20 columns):
id                26753 non-null int64
popularity        26753 non-null float64
budget            26753 non-null int64
revenue           26753 non-null int64
original_title    26753 non-null object
cast              26753 non-null object
director          26753 non-null object
keywords          23523 non-null object
runtime           26753 non-null int64
production_companies 24650 non-null object
vote_count        26753 non-null int64
vote_average      26753 non-null float64
release_year      26753 non-null int64
budget_adj        26753 non-null float64
revenue_adj       26753 non-null float64
profit            26753 non-null int64
profit_adj        26753 non-null float64
profit_ratio      26753 non-null float64
profit_ratio_adj  26753 non-null float64
genre_split       26753 non-null object
dtypes: float64(7), int64(7), object(6)
memory usage: 4.3+ MB
```

- 'keywords' and 'production_companies' have null value but not affect the analysis that not related to them.

In [19]: `df_split_genre.shape`

Out[19]: (26753, 20)

In [20]: `df_split_genre.describe()`

Out[20]:

	id	popularity	budget	revenue	runtime	vote_count	vc
count	26753.000000	26753.000000	2.675300e+04	2.675300e+04	26753.000000	26753.000000	26
mean	58236.098045	0.710244	1.763665e+07	4.779885e+07	103.048892	251.691436	
std	86350.207583	1.118093	3.470727e+07	1.326446e+08	29.560855	640.123565	
min	5.000000	0.000188	0.000000e+00	0.000000e+00	0.000000	10.000000	
25%	10184.000000	0.226575	0.000000e+00	0.000000e+00	90.000000	18.000000	
50%	18065.000000	0.414311	2.500000e+04	0.000000e+00	100.000000	44.000000	
75%	57718.000000	0.779596	2.000000e+07	3.132790e+07	112.000000	176.000000	
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	

Split up 'keywords' columns

```
In [21]: df_split_keywords = df_keywords.copy()
split_keywords = df_split_keywords['keywords'].str.split('|').apply(pd.Series,
1).stack().reset_index(level=1, drop=True)
split_keywords.name = 'keywords_split'
df_split_keywords = df_split_keywords.drop(['keywords'], axis=1).join(split_ke
ywords)
df_split_keywords.head(3)
```

Out[21]:

	id	popularity	budget	revenue	original_title	cast	director	runtime
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	124
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	124
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	124

In [22]: `df_split_keywords.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 37235 entries, 0 to 10865
Data columns (total 20 columns):
id                37235 non-null int64
popularity        37235 non-null float64
budget            37235 non-null int64
revenue           37235 non-null int64
original_title    37235 non-null object
cast              37235 non-null object
director          37235 non-null object
runtime           37235 non-null int64
genres            37235 non-null object
production_companies 35234 non-null object
vote_count        37235 non-null int64
vote_average      37235 non-null float64
release_year      37235 non-null int64
budget_adj        37235 non-null float64
revenue_adj       37235 non-null float64
profit            37235 non-null int64
profit_adj        37235 non-null float64
profit_ratio      37235 non-null float64
profit_ratio_adj  37235 non-null float64
keywords_split    37235 non-null object
dtypes: float64(7), int64(7), object(6)
memory usage: 6.0+ MB
```

- 'production_companies' has null value but not affect the analysis that not related to it.

In [23]: `df_split_keywords.shape`

Out[23]: (37235, 20)

In [24]: `df_split_keywords.describe()`

Out[24]:

	id	popularity	budget	revenue	runtime	vote_count	vc
count	37235.000000	37235.000000	3.723500e+04	3.723500e+04	37235.000000	37235.000000	37
mean	52963.982463	0.783637	1.879837e+07	5.387756e+07	104.437787	289.243776	
std	83038.898888	1.159073	3.494342e+07	1.364298e+08	27.254729	675.433227	
min	5.000000	0.000188	0.000000e+00	0.000000e+00	0.000000	10.000000	
25%	9562.000000	0.255858	0.000000e+00	0.000000e+00	91.000000	21.000000	
50%	14459.000000	0.467556	2.200000e+06	7.707060e+05	101.000000	60.000000	
75%	49010.000000	0.890557	2.360000e+07	4.400869e+07	114.000000	227.000000	
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	

Split up 'production_companies' columns

```
In [25]: df_split_production = df_production.copy()
split_production = df_split_production['production_companies'].str.split('|').
apply(pd.Series,1).stack().reset_index(level=1, drop=True)
split_production.name = 'production_split'
df_split_production = df_split_production.drop(['production_companies'], axis=
1).join(split_production)
df_split_production.head(3)
```

Out[25]:

	id	popularity	budget	revenue	original_title	cast	director	
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	monster d rex v
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	monster d rex v
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	monster d rex v

In [26]: `df_split_production.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 23143 entries, 0 to 10865
Data columns (total 20 columns):
id                23143 non-null int64
popularity        23143 non-null float64
budget            23143 non-null int64
revenue           23143 non-null int64
original_title    23143 non-null object
cast              23143 non-null object
director          23143 non-null object
keywords          20804 non-null object
runtime           23143 non-null int64
genres            23143 non-null object
vote_count        23143 non-null int64
vote_average      23143 non-null float64
release_year      23143 non-null int64
budget_adj        23143 non-null float64
revenue_adj       23143 non-null float64
profit            23143 non-null int64
profit_adj        23143 non-null float64
profit_ratio      23143 non-null float64
profit_ratio_adj  23143 non-null float64
production_split  23143 non-null object
dtypes: float64(7), int64(7), object(6)
memory usage: 3.7+ MB
```

- 'keywords' has null value but not affect the analysis that not related to it.

In [27]: `df_split_production.shape`

Out[27]: (23143, 20)

In [28]: `df_split_production.describe()`

Out[28]:

	id	popularity	budget	revenue	runtime	vote_count	vc
count	23143.000000	23143.000000	2.314300e+04	2.314300e+04	23143.000000	23143.000000	23
mean	63782.611546	0.816030	2.084325e+07	5.533152e+07	104.915482	308.415158	
std	90008.832896	1.209748	3.624897e+07	1.365742e+08	25.835334	681.717341	
min	5.000000	0.000188	0.000000e+00	0.000000e+00	0.000000	10.000000	
25%	9905.000000	0.267255	0.000000e+00	0.000000e+00	92.000000	22.000000	
50%	18228.000000	0.484139	4.000000e+06	7.918300e+05	101.000000	65.000000	
75%	74751.500000	0.937272	2.600000e+07	4.819070e+07	114.000000	257.000000	
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	877.000000	9767.000000	

Split up 'cast' columns

```
In [29]: df_split_cast = df.copy()
split_cast = df_split_cast['cast'].str.split('|').apply(pd.Series,1).stack().r
reset_index(level=1, drop=True)
split_cast.name = 'cast_split'
df_split_cast = df_split_cast.drop(['cast'], axis=1).join(split_cast)
df_split_cast.head(3)
```

Out[29]:

	id	popularity	budget	revenue	original_title	director	keyword
0	135397	32.985763	150000000	1513528810	Jurassic World	Colin Trevorrow	monster dna tyrannosaurus rex velociraptor island
0	135397	32.985763	150000000	1513528810	Jurassic World	Colin Trevorrow	monster dna tyrannosaurus rex velociraptor island
0	135397	32.985763	150000000	1513528810	Jurassic World	Colin Trevorrow	monster dna tyrannosaurus rex velociraptor island

In [30]: df_split_cast.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 52334 entries, 0 to 10865
Data columns (total 20 columns):
id                52334 non-null int64
popularity        52334 non-null float64
budget            52334 non-null int64
revenue           52334 non-null int64
original_title    52334 non-null object
director          52334 non-null object
keywords          45600 non-null object
runtime           52334 non-null int64
genres            52334 non-null object
production_companies 48020 non-null object
vote_count        52334 non-null int64
vote_average      52334 non-null float64
release_year      52334 non-null int64
budget_adj        52334 non-null float64
revenue_adj       52334 non-null float64
profit            52334 non-null int64
profit_adj        52334 non-null float64
profit_ratio      52334 non-null float64
profit_ratio_adj  52334 non-null float64
cast_split        52334 non-null object
dtypes: float64(7), int64(7), object(6)
memory usage: 8.4+ MB
```

- 'keywords' and 'production_companies' have null value but not affect the analysis that not related to them.

In [31]: `df_split_cast.shape`

Out[31]: (52334, 20)

In [32]: `df_split_cast.describe()`

Out[32]:

	id	popularity	budget	revenue	runtime	vote_count	vc
count	52334.000000	52334.000000	5.233400e+04	5.233400e+04	52334.000000	52334.000000	52
mean	63949.016949	0.663357	1.516689e+07	4.132668e+07	102.954618	224.558490	
std	90563.719104	1.013682	3.136560e+07	1.189496e+08	28.888852	585.220224	
min	5.000000	0.000188	0.000000e+00	0.000000e+00	0.000000	10.000000	
25%	10448.000000	0.216906	0.000000e+00	0.000000e+00	90.000000	17.000000	
50%	19621.000000	0.394209	0.000000e+00	0.000000e+00	99.000000	40.000000	
75%	71866.000000	0.733947	1.700000e+07	2.681011e+07	112.000000	154.000000	
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	

Split up 'director' columns

```
In [33]: df_split_director = df.copy()
split_director = df_split_director['director'].str.split('|').apply(pd.Series,
1).stack().reset_index(level=1, drop=True)
split_director.name = 'director_split'
df_split_director = df_split_director.drop(['director'], axis=1).join(split_director)
df_split_director.head(3)
```

Out[33]:

	id	popularity	budget	revenue	original_title	cast	
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	monster c rex
1	76341	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	apocalyptic
2	262500	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	novel revolution dyst

In [34]: `df_split_director.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 11774 entries, 0 to 10865
Data columns (total 20 columns):
id                11774 non-null int64
popularity        11774 non-null float64
budget            11774 non-null int64
revenue           11774 non-null int64
original_title    11774 non-null object
cast              11774 non-null object
keywords          10209 non-null object
runtime           11774 non-null int64
genres            11774 non-null object
production_companies 10708 non-null object
vote_count        11774 non-null int64
vote_average      11774 non-null float64
release_year      11774 non-null int64
budget_adj        11774 non-null float64
revenue_adj       11774 non-null float64
profit            11774 non-null int64
profit_adj        11774 non-null float64
profit_ratio      11774 non-null float64
profit_ratio_adj  11774 non-null float64
director_split    11774 non-null object
dtypes: float64(7), int64(7), object(6)
memory usage: 1.9+ MB
```

- 'keywords' and 'production_companies' have null value but not affect the analysis that not related to them.

In [35]: `df_split_director.shape`

Out[35]: (11774, 20)

In [36]: `df_split_director.describe()`

Out[36]:

	id	popularity	budget	revenue	runtime	vote_count	vc
count	11774.000000	11774.000000	1.177400e+04	1.177400e+04	11774.000000	11774.000000	11
mean	67035.732631	0.655070	1.478524e+07	4.080166e+07	103.047138	221.918379	
std	92428.824638	1.005885	3.134590e+07	1.195286e+08	41.075401	580.822606	
min	5.000000	0.000188	0.000000e+00	0.000000e+00	0.000000	10.000000	
25%	10705.250000	0.209859	0.000000e+00	0.000000e+00	90.000000	17.000000	
50%	21330.000000	0.386192	0.000000e+00	0.000000e+00	98.000000	39.000000	
75%	78382.500000	0.722796	1.500000e+07	2.417932e+07	111.000000	149.000000	
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	

- 'keywords' and 'production_companies' have null value but not affect the analysis that not related to them.

Now we have 5 clean dataframes:

- df
- df_keywords
- df_production
- df_cast
- df_director

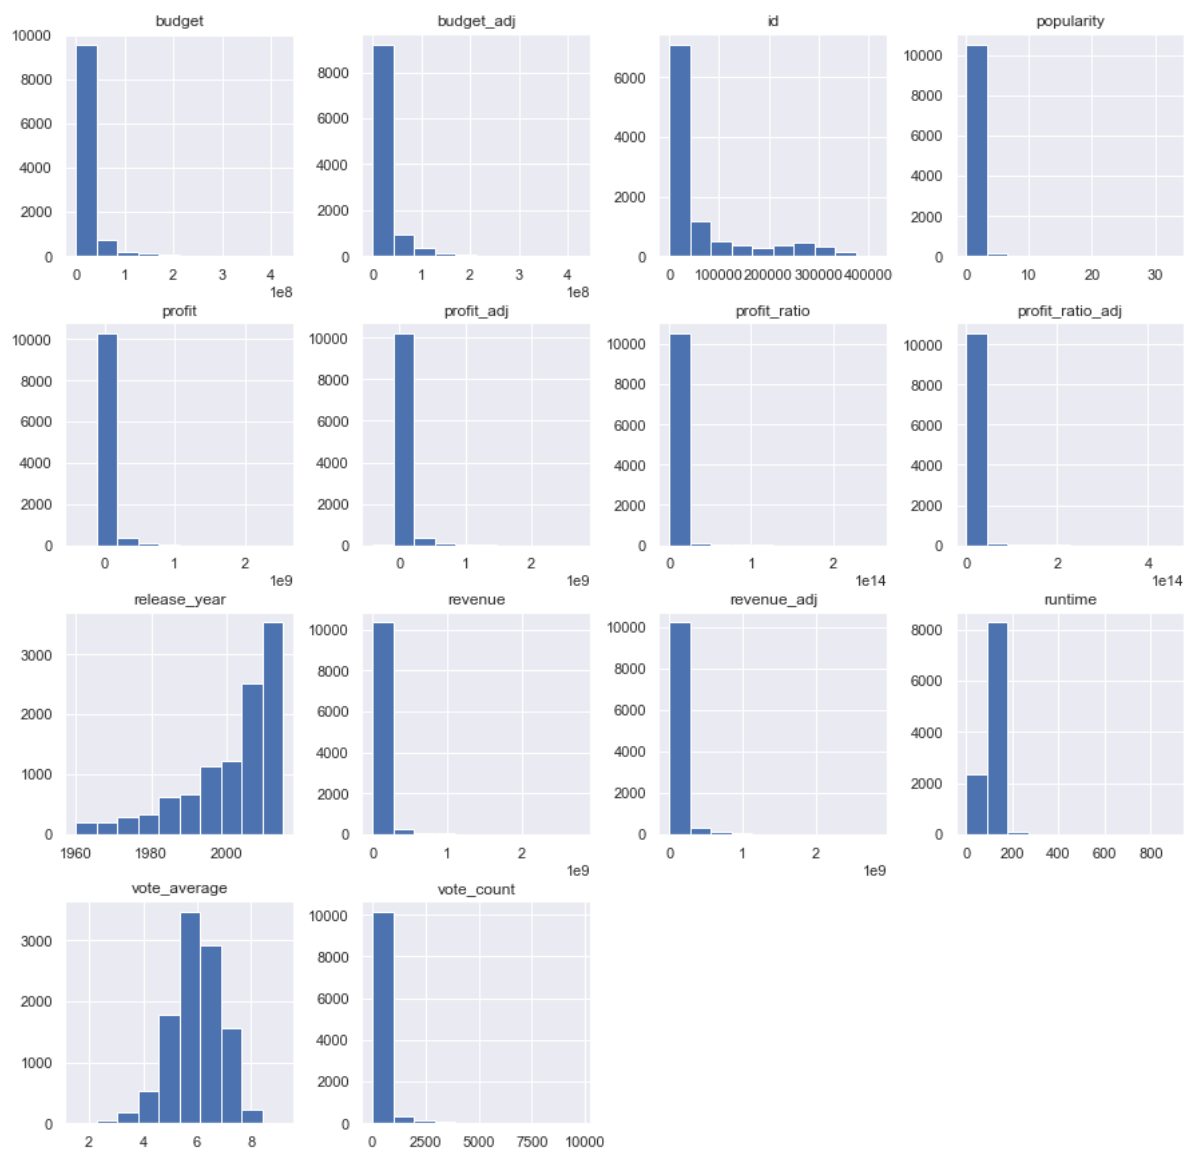
Exploratory Data Analysis

In [37]: `df.corr()`


Out[37]:

	id	popularity	budget	revenue	runtime	vote_count	vote_average
id	1.000000	-0.009464	-0.138935	-0.097424	-0.083996	-0.032767	-0.071896
popularity	-0.009464	1.000000	0.544240	0.662843	0.138278	0.800619	0.217906
budget	-0.138935	0.544240	1.000000	0.734487	0.192168	0.632074	0.087318
revenue	-0.097424	0.662843	0.734487	1.000000	0.164276	0.790889	0.178477
runtime	-0.083996	0.138278	0.192168	0.164276	1.000000	0.164966	0.177276
vote_count	-0.032767	0.800619	0.632074	0.790889	0.164966	1.000000	0.260554
vote_average	-0.071896	0.217906	0.087318	0.178477	0.177276	0.260554	1.000000
release_year	0.510385	0.093044	0.119004	0.059072	-0.119286	0.110317	-0.127746
budget_adj	-0.186980	0.512098	0.968881	0.705949	0.222645	0.586298	0.099925
revenue_adj	-0.137099	0.608384	0.621809	0.918990	0.177397	0.707517	0.199418
profit	-0.073557	0.628699	0.569730	0.976173	0.137497	0.755681	0.188220
profit_adj	-0.107138	0.562359	0.452854	0.867967	0.143911	0.656523	0.202922
profit_ratio	-0.058220	-0.004752	-0.075339	0.036690	0.038097	-0.023790	0.022227
profit_ratio_adj	-0.070609	-0.013727	-0.073365	0.030528	0.037407	-0.029740	0.029567

```
In [38]: df.hist(figsize=(15,15));
```



```
In [39]: # Scatter plot of df using seaborn
sns_scatterplt = sns.pairplot(df, height = 2.5)
```



From scatter plots:

- Revenue and Profits are positively correlated

Research Question 0: Fun Facts

This section I will investigate various fun facts

- 1. Words with highest frequency appeared in the movie titles:

```
In [40]: # Wordcloud for title visualization:
```

```
from wordcloud import WordCloud, STOPWORDS
text = (str(df['original_title']))
plt.subplots(figsize=(15,15))
wordcloud = WordCloud(stopwords=STOPWORDS, background_color='white', width=1500, height=1200).generate(text)
plt.imshow(wordcloud)
plt.title('Title')
plt.axis('off');
```



- **Title analysis just for fun. ^^**

- **2. Top 10 Most Popular Movies:**

In [41]: *# Top 10 Most Popular Movies:*

```
df[['popularity', 'original_title']].sort_values(by='popularity', ascending=False).head(10)
```

Out[41]:

	popularity	original_title
0	32.985763	Jurassic World
1	28.419936	Mad Max: Fury Road
629	24.949134	Interstellar
630	14.311205	Guardians of the Galaxy
2	13.112507	Insurgent
631	12.971027	Captain America: The Winter Soldier
1329	12.037933	Star Wars
632	11.422751	John Wick
3	11.173104	Star Wars: The Force Awakens
633	10.739009	The Hunger Games: Mockingjay - Part 1

- 3. Top 10 Highest Rating Movies:

In [42]: *# Top 10 Highest Rating Movies:*

```
df[['vote_average', 'original_title']].sort_values(by='vote_average', ascending=False).head(10)
```

Out[42]:

	vote_average	original_title
3894	9.2	The Story of Film: An Odyssey
1200	8.8	Black Mirror: White Christmas
6911	8.7	Pink Floyd: Pulse
3690	8.5	The Art of Flight
8221	8.5	A Personal Journey with Martin Scorsese Throug...
8839	8.5	Dave Chappelle: Killin' Them Softly
8411	8.5	Queen - Rock Montreal
4178	8.4	The Shawshank Redemption
2334	8.4	Rush: Beyond the Lighted Stage
609	8.4	The Jinx: The Life and Deaths of Robert Durst

- 4. Top 10 Most Profitable Movies (sorted by adjusted profit and profit):

```
In [43]: df[['profit_adj', 'original_title']].sort_values(by='profit_adj', ascending=False).head(10)
```

Out[43]:

	profit_adj	original_title
1329	2.750137e+09	Star Wars
1386	2.586237e+09	Avatar
5231	2.234714e+09	Titanic
10594	2.128036e+09	The Exorcist
9806	1.878643e+09	Jaws
8889	1.767968e+09	E.T. the Extra-Terrestrial
3	1.718723e+09	Star Wars: The Force Awakens
8094	1.551568e+09	The Net
10110	1.545635e+09	One Hundred and One Dalmatians
7309	1.376998e+09	The Empire Strikes Back

```
In [44]: df[['profit', 'original_title']].sort_values(by='profit', ascending=False).head(10)
```

Out[44]:

	profit	original_title
1386	2544505847	Avatar
3	1868178225	Star Wars: The Force Awakens
5231	1645034188	Titanic
0	1363528810	Jurassic World
4	1316249360	Furious 7
4361	1299557910	The Avengers
3374	1202817822	Harry Potter and the Deathly Hallows: Part 2
14	1125035767	Avengers: Age of Ultron
5422	1124219009	Frozen
8094	1084279658	The Net

Profitability should be evaluated by profit rather than ratios. Some movies have "0" budget so we add 0.000001 in previous steps, and this will make profit_ratio super high.

- 5. Top 10 Actors Starred In Most Movies:

```
In [45]: df_split_cast['cast_split'].value_counts().head(10)
```

```
Out[45]: Robert De Niro      72
         Samuel L. Jackson  71
         Bruce Willis      62
         Nicolas Cage       61
         Michael Caine     53
         Robin Williams    51
         John Cusack       50
         Morgan Freeman    49
         John Goodman      49
         Liam Neeson       48
         Name: cast_split, dtype: int64
```

• 6. Top 10 Keywords In Most Movies:

```
In [46]: df_split_keywords['keywords_split'].value_counts().head(10)
```

```
Out[46]: woman director      408
         independent film    393
         based on novel     278
         sex                 272
         sport               215
         murder              204
         musical             169
         biography           168
         new york            162
         suspense            159
         Name: keywords_split, dtype: int64
```

• 7. Top 10 Production Companies:


```
In [47]: df_split_production['production_split'].value_counts().head(10)
```

```
Out[47]: Universal Pictures      522
         Warner Bros.           509
         Paramount Pictures      431
         Twentieth Century Fox Film Corporation  282
         Columbia Pictures       272
         New Line Cinema         219
         Metro-Goldwyn-Mayer (MGM)  218
         Walt Disney Pictures     213
         Touchstone Pictures      178
         Columbia Pictures Corporation  160
         Name: production_split, dtype: int64
```


Research Question 1: Genre Trends from 1960 to 2015

This section I will investigate different genres

```
In [48]: # Plot pie chart to visualize genre distribution
df_split_genre['genre_split'].value_counts().plot(kind = 'pie',figsize = (16,16));
plt.title('Genre of Movies, 1960-2015', size=30)
plt.ylabel('Genre', size=20);
```



```
In [49]: #Plot bar chart to visualize genre distribution
df_split_genre['genre_split'].value_counts().plot(kind='bar',figsize = (12,12
));
plt.title('Genre of Movies, 1960-2015', size=30)
plt.xlabel('Genre', size=20)
plt.ylabel('Movie #', size=20);
```



```
In [50]: # Select data from df for 5 most popular genres: Drama, Comedy, Thriller, Action, Romance.
# Then plot the total counts of different genres for each year from 1960 to 2015
drama = df_split_genre.genre_split == 'Drama'
df_drama = df_split_genre[drama]
df_drama.groupby('release_year')['genre_split'].count().plot(figsize=(15,15),label='Drama')

comedy = df_split_genre.genre_split == 'Comedy'
df_comedy = df_split_genre[comedy]
df_comedy.groupby('release_year')['genre_split'].count().plot(label='Comedy')


thriller = df_split_genre.genre_split == 'Thriller'
df_thriller = df_split_genre[thriller]
df_thriller.groupby('release_year')['genre_split'].count().plot(label='Thriller')

action = df_split_genre.genre_split == 'Action'
df_action = df_split_genre[action]
df_action.groupby('release_year')['genre_split'].count().plot(label='Action')

romance = df_split_genre.genre_split == 'Romance'
df_romance = df_split_genre[romance]
df_romance.groupby('release_year')['genre_split'].count().plot(label='Romance')

plt.title('5 Most Popular Movie Genres Trendline from 1960 to 2015',size=30)
plt.xlabel('Year',size=20)
plt.ylabel('# of Movies',size=20)
plt.legend(fontsize = 'xx-large');
```

5 Most Popular Movie Genres Trendline from 1960 to 2015




```
In [51]: # Select data from df for 5 most popular genres: Drama, Comedy, Thriller, Action, Romance.
# Then plot the total counts of different genres for each year from 1960 to 2015

top_genres = ['Drama', 'Comedy', 'Thriller', 'Action', 'Romance']

##### Code below is hard to read but more concise, I don't know which one is better according to rubric #####
for i in range(len(top_genres)):
    df_split_genre[df_split_genre.genre_split == top_genres[i]].groupby('release_year')['genre_split'].count().plot(figsize=(15,15),label=top_genres[i])

plt.title('5 Most Popular Movie Genres Trendline from 1960 to 2015',size=30)
plt.xlabel('Year',size=20)
plt.ylabel('# of Movies',size=20)
plt.legend(fontsize = 'xx-large');
```



From the line chart above:

- All types of movies are increasing from 1960 to 2015
- The largest growth rate occurred during 2000-2010
- Drama is the most popular genre through the years except being exceeded by Comedy during the late 80's
- Romance is the least popular genre among top 5 genres. The growth rate of romance movies is the slowest.

Research Question 2: Properties Associated With Higher Profits

This section I will investigate which movie factors are related to higher profits

• 1. Profit vs. Genre:

In [52]: *# Find all genres:*

```
df_split_genre.genre_split.unique()
```

Out[52]: array(['Action', 'Adventure', 'Science Fiction', 'Thriller', 'Fantasy',
 'Crime', 'Western', 'Drama', 'Family', 'Animation', 'Comedy',
 'Mystery', 'Romance', 'War', 'History', 'Music', 'Horror',
 'Documentary', 'TV Movie', 'Foreign'], dtype=object)

In [53]: *# Get total number of genres for later avg_profit calculation:*

```
num_g = df_split_genre.genre_split.value_counts()
genres_total_num = list(num_g)
genres_total_num
```

Out[53]: [4746,
3775,
2902,
2376,
1708,
1636,
1465,
1353,
1221,
1214,
908,
808,
664,
470,
399,
330,
268,
184,
164,
162]

```
In [54]: # Calculate sums of profit of each genres:

genres_labels = ['Action', 'Adventure', 'Science Fiction', 'Thriller', 'Fantasy',
                 'Crime', 'Western', 'Drama', 'Family', 'Animation', 'Comedy',
                 'Mystery', 'Romance', 'War', 'History', 'Music', 'Horror',
                 'Documentary', 'TV Movie', 'Foreign']

profit_by_genres = []


for i in range(len(genres_labels)):
    profit_by_genres.append(df_split_genre[df_split_genre.genre_split == genres_labels[i]].profit.sum())

profit_by_genres
```

```
Out[54]: [107439517424,
          111199018978,
          55511321460,
          71284730705,
          58355181708,
          33450381145,
          1583109216,
          82594648101,
          60420445751,
          36417750351,
          91896372240,
          19519620245,
          39644299221,
          7212590243,
          4488235887,
          7813519034,
          17346109400,
          745802029,
          -2700000,
          9406683]
```


In [55]: *# Visualization for total profits of each genres:*

```
plt.figure(figsize=(16, 16))
sns.barplot(x=genres_labels, y=profit_by_genres)
plt.title('Total Profits for each Genres from 1960 to 2015',size=30)
plt.xlabel('Genres',size=20)
plt.ylabel('Total Profit ($)',size=20)
plt.xticks(fontsize=20,rotation=90)
plt.yticks(fontsize=20);
```




In [56]: *# Calculate average profits for each genres:*

```
avg_profit = np.divide(profit_by_genres,genres_total_num)
avg_profit = avg_profit.astype(int)
avg_profit
```

Out[56]: array([22637909, 29456693, 19128642, 30001991, 34165797, 20446443,
 1080620, 61045564, 49484394, 29998146, 101207458, 24157945,
 59705269, 15345936, 11248711, 23677330, 64724288, 4053271,
 -16463, 58065])

In [57]: *# Visualization for average profits of each genres:*

```
plt.figure(figsize=(16, 16))
sns.barplot(x=genres_labels, y=avg_profit)
plt.title('Average Profits for each Genres from 1960 to 2015',size=30)
plt.xlabel('Genres',size=20)
plt.ylabel('Average Profit ($)',size=20)
plt.xticks(fontsize=20,rotation=90)
plt.yticks(fontsize=20);
```




From bar chart above:

- Top 5 most profitable genres are: 1.Comedy, 2.Horror, 3.Drama, 4.Romance, 5.Family.

- 2. Profit vs. Popularity:

```
In [58]: # Plot Profit vs. Popularity:


plt.figure(figsize=(10, 10))
sns.scatterplot(x=df['popularity'],y=df['profit'])
plt.title('Profits vs. Popularity',size=30)
plt.xlabel('Popularity',size=20)
plt.ylabel('Profit ($)',size=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15);
```



- 3. Profit vs. Runtime:

In [59]: *# Plot Profit vs. Runtime:*


```
plt.figure(figsize=(10, 10))
sns.scatterplot(x=df['runtime'],y=df['profit'])
plt.title('Profits vs. Runtime',size=30)
plt.xlabel('Runtime (min)',size=20)
plt.ylabel('Profit ($)',size=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15);
```



- **4. Profit vs. Revenue:**

In [60]: *# Plot Profit vs. Revenue:*


```
plt.figure(figsize=(10, 10))
sns.scatterplot(x=df['revenue'],y=df['profit'])
plt.title('Profits vs. Revenue',size=30)
plt.xlabel('Revenue ($)',size=20)
plt.ylabel('Profit ($)',size=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15);
```



- **5. Profit vs. Budget:**

In [61]: *# Plot Profit vs. Budget:*


```
plt.figure(figsize=(10, 10))
sns.scatterplot(x=df['budget'],y=df['profit'])
plt.title('Profits vs. Budget',size=30)
plt.xlabel('Budget ($)',size=20)
plt.ylabel('Profit ($)',size=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15);
```



- **6. Profit vs. Vote_Count:**

In [62]: *# Plot Profit vs. Vote_Count:*

```
plt.figure(figsize=(10, 10))
sns.scatterplot(x=df['vote_count'],y=df['profit'])
plt.title('Profits vs. Vote Count',size=30)
plt.xlabel('Vote Count',size=20)
plt.ylabel('Profit ($)',size=20)
plt.xticks(fontsize=15)
plt.yticks(fontsize=15);
```



Conclusions

Summary of Data

Question 0: Fun Facts

- Top 5 most common words in movie titles: *Max, Man, Big, Carry, Hand.*
- Top 10 Most Popular Movies: *Jurassic World, Mad Max: Fury Road, Interstellar, Guardians of the Galaxy, Insurgent, Captain America: The Winter Soldier, Star Wars, John Wick, Star Wars: The Force Awakens, The Hunger Games: Mockingjay - Part 1.*
- Top 10 Highest Rating Movies: *The Story of Film: An Odyssey, Black Mirror: White Christmas, Pink Floyd: Pulse, The Art of Flight, A Personal Journey With Martin Scorsese Through American Movies, Dave Chappelle: Killin' Them Softly, Queen - Rock Montreal, The Shawshank Redemption, Rush: Beyond the Lighted Stage, The Jinx: The Life and Deaths of Robert Durst.*
- Top 10 Most Profitable Movies: *Avatar, Star Wars: The Force Awakens, Titanic, Jurassic World, Furious 7, The Avengers, Harry Potter and the Deathly Hallows: Part 2, Avengers: Age of Ultron, Frozen, The Net.*
- Top 10 Most Profitable Movies (sorted by adjusted profit): *Star Wars, Avatar, Titanic, The Exorcist, Jaws, E.T. the Extra-Terrestrial, Star Wars: The Force Awakens, The Net, One Hundred and One Dalmatians, The Empire Strikes Back.*
- Top 10 Actors Starred Most Movies: *Robert De Niro, Samuel L. Jackson, Bruce Willis, Nicolas Cage, Michael Caine, Robin Williams, John Cusack, Morgan Freeman, John Goodman, Susan Sarandon.*
- Top 10 Keywords: *woman director, independent film, based on novel, sex, sport, murder, musical, biography, new york, suspense.*
- Top 10 Production Companies: *Universal Pictures, Warner Bros., Paramount Pictures, Twentieth Century Fox Film Corporation, Columbia Pictures, New Line Cinema, Metro-Goldwyn-Mayer (MGM), Walt Disney Pictures, Touchstone Pictures, Columbia Pictures Corporation.*

NOTE: I'm not sure if *Columbia Pictures* and *Columbia Pictures Corporation* are the same company. If so, I should replace one name to the other to merge those two.

Question 1: Genre Trends from 1960 to 2015

- Drama, Comedy, Thriller, Action and Romance are the most popular genres and make up over 50% of all movies made from 1960-2015. TV Movie, Western, and Foreign are the least popular.
- All types of movies are increasing from 1960 to 2015.
- The largest growth rate occurred during 2000-2010.
- Drama is the most popular genre through the years except being exceeded by Comedy during the late 80's.
- Romance is the least popular genre among top 5 genres. The growth rate of romance movies is the slowest.

Question 2: Properties Associated with Higher Profits

- Top 5 Genres with Highest Total Profits: *Adventure, Action, Comedy, Drama, Thriller.*
- Top 5 Genres with Highest Average Profits: *Comedy, Horror, Drama, Romance, Family.*
- Profit vs. Popularity: Correlation is 0.628699, moderate positive correlation. Higher popularity can somewhat lead to higher profit for a movie.

- **Profit vs. Runtime:** Correlation is 0.137497, weak positive correlation. Runtime is not related to profit for a movie.
- **Profit vs. Revenue:** Correlation is 0.976173, strong positive correlation which is obvious.
- **Profit vs. Budget:** Correlation is 0.569730, moderate positive correlation. Higher budget cannot guarantee higher profit for a movie.
- **Profit vs. Vote Count:** Correlation is 0.755681, medium-strong positive correlation. Since correlation between vote count and popularity is 0.800619 and has strong positive correlation, more vote counts indicate the movie will have higher profit.

Notes & Limitations

Raw Data:

Original data was collected only from *The Movie Database (TMDb)*, so sample bias may exist for data such as 'popularity', 'vote_count' and 'vote_average'. For more accurate results, data from other movie database (eg. IMDb) should be joined into our dataframe.

Data Cleaning:

Columns dropped and reasons:

'imdb_id' : Already have 'id'
'homepage' : Not relevant
'tagline' : Already have keywords
'overview' : Already have keywords
'release_year' : Already have 'release_date'

Added profit = (revenue - budget) and profit_ratio = (profit/budget) columns to investigate profitability.

Split columns 'cast', 'director', 'keywords', 'genres', 'production_companies' that contain multiple values separated by pipe (|).

For the purpose of preserving more data, I did not remove rows with null values in 'keywords' and/or 'production_companies' columns. After Data Cleaning, I have 5 clean dataframes:

- df
- df_keywords
- df_production
- df_cast
- df_director

All of them have null value in either 'keywords' or 'production_companies' or both columns, but this will not affect our analysis because each dataframe will only be using to answer specific questions that not related to the column(s) with null values.